Homework #2: Machine Learning

Kyler Little

February 2, 2018

Problem #1

(a) More generally, if we are learning from ± 1 data to predict a noisy target $P(y|\boldsymbol{x})$ with candidate hypothesis h, show that the maximum likelihood method reduces to the task of finding h that minimizes:

$$E_{\text{in}}(\boldsymbol{w}) = \sum_{n=1}^{N} [\![y_n = +1]\!] \ln(\frac{1}{h(\boldsymbol{x}_n)}) + [\![y_n = -1]\!] \ln \frac{1}{1 - h(\boldsymbol{x}_n)}$$

Formally, we are trying to learn the target function:

$$f(\boldsymbol{x}) = P[y = +1|\boldsymbol{x}]$$

In non-mathematical terms, this target function tries to determine the probability of classifying the data ($\boldsymbol{x}$) as $+1$. The function f is generated by a noisy target function, so the data doesn't give us the value of f explicitly. Instead, we try to minimize the error between the classifications of actual training data and the predicted classifications from our model. We can do this by talking about the notion of likelihood. The likelihood that the target distribution $P(y|\boldsymbol{x})$ is captured by our hypothesis $h(\boldsymbol{x})$ is:

$$P(y|\boldsymbol{x}) = \begin{cases} h(\boldsymbol{x}) & \text{for } y = +1; \\ 1 - h(\boldsymbol{x}) & \text{for } y = -1. \end{cases}$$

If we assume that the data points were independently generated (a very fair assumption to make), then we can express the probability of getting all of the y_n's in the data set from their corresponding x_n's as:

$$\prod_{n=1}^{N} P(y_n|\boldsymbol{x}_n)$$

The method of maximum likelihood method would selects the hypothesis h which maximizes the probability above. We can actually express this maximization problem as a minimization problem by using the properties of ln. The natural logarithm is a monotonically increasing function, and so taking the natural logarithm of the probability above will not affect the maximization problem. By the same token, $-\ln$ is a monotonically decreasing function, so we can take the natural logarithm of the probability above and convert the maximization problem to a minimization problem without affecting global minima of the probability distribution. Using a simple logarithm rule ($-\ln x = \ln x^{-1}$), we find:

$$\prod_{n=1}^{N} P(y_n|\boldsymbol{x}_n) = \sum_{n=1}^{N} \ln(\frac{1}{P(y_n|\boldsymbol{x}_n)})$$

Substituting in what $P(y|\boldsymbol{x})$ actually equals, we arrive at what we wanted. We now find h that minimizes:

$$E_{\text{in}}(\boldsymbol{w}) = \sum_{n=1}^{N} [\![y_n = +1]\!] \ln(\frac{1}{h(\boldsymbol{x}_n)}) + [\![y_n = -1]\!] \ln \frac{1}{1 - h(\boldsymbol{x}_n)}$$

(b) For the case $h(\boldsymbol{x}) = \theta(\boldsymbol{w}^T\boldsymbol{x})$, argue that minimizing the in-sample error in part (a) is equivalent to minimizing the one in (3.9).
For two probabilistic distributions $p, 1-p$ and $q, 1-q$ with binary outcomes, the cross-entropy (from information theory) is:

$$p \log \frac{1}{q} + (1-p) \log \frac{1}{1-q}$$

The in-sample error in part (a) corresponds to a cross-entropy error measure on the data point $(\boldsymbol{x}_n, y_n)$, with $p = [\![y_n = +1]\!]$ and $q = h(\boldsymbol{x}_n)$.

Problem #2

Recall the objective function for linear regression can be expressed as

$$E(w) = \frac{1}{N}||Xw - y||^2$$

as in Equation (3.3) of LFD. Minimizing this function with respect to w leads to the optimal w as $(XTX)^{-1}X^Ty$. This solution holds only when X^TX is

nonsingular. To overcome this problem, the following objective function is commonly minimized instead:

$$E_2(w) = ||Xw - y||^2 + \lambda||w||^2$$

where $\lambda > 0$ is a user-specified parameter. Please do the following:
(a) Derive the optimal w that minimize $E_2(w)$.
(b) Explain how this new objective function can overcome the singularity problem of X^TX.

Problem #3

In logistic regression, the objective function can be written as:

$$E(w) = \frac{1}{N}\sum_{n=1}^{N} \ln(1 + e^{-y_n w^T x_n})$$

(a) (10 points) Compute the first-order derivative $\nabla E(w)$. You will need to provide the intermediate steps of derivation.
(b) (10 points) Once the optimal w is obtained, it will be used to make predictions as follows:

$$\text{Predicted class of x} = \begin{cases} 1 & \text{if } \theta(w^T x) \geq 0.5 \\ -1 & \text{if } \theta(w^T x) < 0.5 \end{cases}$$

where the function $\theta(z) = \frac{1}{1+e^{-z}}$ looks like

Explain why the decision boundary of logistic regression is still linear, though the linear signal w^Tx is passed through a nonlinear function θ to compute the outcome of prediction. Is the decision boundary still linear if the prediction rule is changed to the following? Justify briefly.

$$\text{Predicted class of x} = \begin{cases} 1 & \text{if } \theta(w^T x) \geq 0.9 \\ -1 & \text{if } \theta(w^T x) < 0.9 \end{cases}$$

In light of your answers to the above two questions, what is the essential property of logistic regression that results in the linear decision boundary?

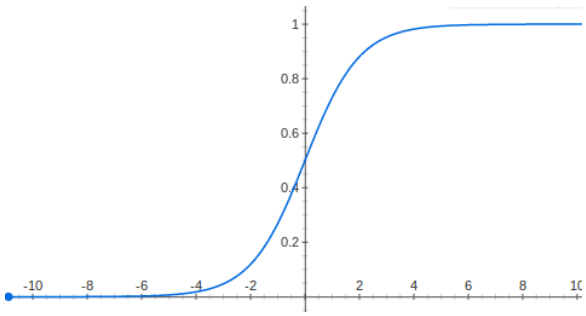

Figure 1: A basic sigmoid curve

Problem #4

As your final deliverable to a customer, would you use the linear model with or without the 3rd order polynomial transform? Briefly explain your reasoning.