Homework #3: Machine Learning

Kyler Little

February 27, 2018

Problem #1

Suppose that we have three coloured boxes r (red), b (blue), and g (green). Box r contains 3 apples, 4 oranges, and 3 limes; box b contains 1 apple, 1 orange, and 0 limes; and box g contains 3 apples, 3 oranges, and 4 limes. If a box is chosen at random with probabilities $p(r) = 0.2$, $p(b) = 0.2$, $p(g) = 0.6$, and a piece of fruit is removed from the box (with equal probability of selecting any of the items in the box), then what is the probability of selecting an apple? If we observe that the selected fruit is in fact an orange, what is the probability that it came from the green box?

In regards to the first question, we need to compute the marginal probability $p(\text{apple})$. We can do this using the sum and product rules. Let Y denote coloured box.

$$
\begin{aligned}
p(\text{apple}) &= \sum_Y p(\text{apple}, Y) = \sum_Y p(\text{apple}|Y)p(Y) \\
&= p(\text{apple}|r)p(r) + p(\text{apple}|b)p(b) + p(\text{apple}|g)p(g) \\
&= 0.3(0.2) + 0.5(0.2) + 0.3(0.6) \\
&= 0.34
\end{aligned}
$$

The second question asks us to compute $p(g|\text{orange})$. We have all of the conditional probabilities of the form $P(\text{fruit}|\text{box})$, and we want $P(\text{box}|\text{fruit})$. This is exactly the purpose of Bayes' Theorem, and so we apply it.

$$
\begin{aligned}
p(g|\text{orange}) &= \frac{p(\text{orange}|g) * p(g)}{p(\text{orange})} \\
&= \frac{p(\text{orange}|g) * p(g)}{p(\text{orange}|r)p(r) + p(\text{orange}|b)p(b) + p(\text{orange}|g)p(g)} \\
&= \frac{0.3(0.6)}{0.4(0.2) + 0.5(0.2) + 0.3(0.6)} \\
&= 0.5
\end{aligned}
$$

Problem #2

Given the following data set containing three attributes and one class, use naïve Bayes classifier to determine the class (Yes/No) of Stolen for a Red Domestic SUV.

Example No.	Color	Type	Origin	Stolen?
1	Red	Sports	Domestic	Yes
2	Red	Sports	Domestic	No
3	Red	Sports	Domestic	Yes
4	Yellow	Sports	Domestic	No
5	Yellow	Sports	Imported	Yes
6	Yellow	SUV	Imported	No
7	Yellow	SUV	Imported	Yes
8	Yellow	SUV	Domestic	No
9	Red	SUV	Imported	No
10	Red	Sports	Imported	Yes

Let $x^* = [\text{Red SUV Domestic}]^T$. To use naïve Bayes' classifier, we shall calculate $P(\text{Yes} \,|\, x^*)$ and $P(\text{No} \,|\, x^*)$. Whichever is larger will be the decision. To start, we will calculate $P(\text{Yes} \,|\, x^*)$. We ignore the denominators because they are the same, and we only care which is larger.

$$\begin{aligned} P(\text{Yes} \,|\, x^*) &\propto P(x^* \,|\, \text{Yes}) * P(\text{Yes}) \\ &\propto (P(\text{Red} \,|\, \text{Yes}) * P(\text{SUV} \,|\, \text{Yes}) * P(\text{Domestic} \,|\, \text{Yes})) * P(\text{Yes}) \\ &\propto ((3/5) * (1/5) * (2/5)) * 0.5 \\ &\propto 0.024 \end{aligned}$$

Next, we calculate $P(\text{No} \,|\, x^*)$, again ignoring the denominator.

$$\begin{aligned} P(\text{No} \,|\, x^*) &\propto P(x^* \,|\, \text{No}) * P(\text{No}) \\ &\propto (P(\text{Red} \,|\, \text{No}) * P(\text{SUV} \,|\, \text{No}) * P(\text{Domestic} \,|\, \text{No})) * P(\text{No}) \\ &\propto ((2/5) * (3/5) * (3/5)) * 0.5 \\ &\propto 0.072 \end{aligned}$$

Since $P(\text{No} \,|\, x^*) > P(\text{Yes} \,|\, x^*)$, our decision is "No."

Problem #3

This question is about naïve Bayes classifier. Please do the following:
(a) State what is the simplifying assumption made by naïve Bayes classifier.

The simplifying assumption is that each x_i of the input vector are conditionally independent, given the target value (classification). In other words, the probability of observing $x_1, x_2, ..., x_d$, given the target value, is just the product of the probabilities of the individual attributes: $P(x_1, x_2, ..., x_d \,|\, y) = \prod_{i=1}^{d} P(x_i \,|\, y)$.

(b) Given a binary-class classification problem in which the class labels are binary, the dimension of feature is d, and each attribute can take k different values. Please provide

the numbers of parameters to be estimated with AND without the simplifying assumption. Briefly justify why the simplifying assumption is necessary.

Without the simplifying assumption, one has to enumerate all possible values of the input space, and for each distinct value, calculate its frequency in the training data to obtain an estimate. The problem with this is that the number of these parameters needed to be estimated is the number of target values (i.e. class labels) times the number of distinct values in the input space minus one (the last value need not be estimated since it is 1 minus the sum of the others). This is a very large number of parameters to be estimated. For instance, in our situation, each x_i can take k different values, and there are d attributes. By the multiplication rule for combinatorics, there are $k * k * \cdots * k = k^d$ distinct values of the input space. The last value need not be estimated; in fact, it is not free to vary at all. Thus, the number of parameters to be estimated is the number of degrees of freedom: $k^d - 1$. Since we must estimate this number of parameters for each class or target value, and the class labels are binary, there are $2(k^d - 1)$ parameters to be estimated.

On the other hand, when the simplifying assumption is made, we are able to estimate far less parameters. We need only estimate the number of target values (i.e. class labels) times the number of distinct attributes– $2d$ in this case.

Besides the fact that this simplifying assumption significantly reduces the number of parameters to be estimated ($2(k^d - 1) \ll 2d$, in general), the assumption is necessary for other reasons as well. Without the assumption, the model wouldn't generalize well to unseen data. For instance, if the model hadn't seen a particular input vector in its training data, there wouldn't be a prediction at all, since $P(x \mid y_1) = P(x \mid y_2) = 0$. Thus, an extremely large training data set would be needed in order for the model to function correctly. The size of the training data set N would need to be much larger than k^d, which is very infeasible in reality.

Problem #4

Assume we want to classify science texts into three categories physics, biology and chemistry. The following probabilities have been estimated from analyzing a corpus of pre-classified web-pages gathered from Yahoo.

c	Physics	Biology	Chemistry
P(c)	0.35	0.40	0.25
P(atom | c)	0.1	0.01	0.2
P(carbon | c)	0.005	0.03	0.05
P(proton | c)	0.05	0.001	0.05
P(life | c)	0.001	0.1	0.008
P(earth | c)	0.005	0.006	0.003

Assuming that the probability of each evidence word is independent of other word occurrences given the category of the text, compute the (posterior) probability for each of the possible categories each of the following short texts; and based on that, their most likely classification. Assume that the categories are disjoint and exhaustive (i.e., every text is either physics, or

3

biology or chemistry, and no text can be more than one). Assume that words are first stemmed to reduce them to their base form (atoms → atom) and ignore any words that are not in the table:

A: the carbon atom is the foundation of life on earth
B: the carbon atom contains 12 protons.

This is a job for naïve Bayes' classifier. The first task will be to vectorize A and B. We will use a binary vector, where words 'atom', 'carbon', 'proton', 'life', and 'earth' correspond to entries one through five of the vector. If the word appears at least once, then the entry will be set to 1. Otherwise, it will be 0.

$$A = [1\ 1\ 0\ 1\ 1]^T$$
$$B = [1\ 1\ 1\ 0\ 0]^T$$

The decision will be:

$$\arg\max_v \prod_{j=1}^{5} p_j(x^j \mid y = v)p(y = v)$$

where $v \in \{\text{Physics}, \text{Biology}, \text{Chemistry}\}$. We begin with classifying A.

$$\hat{y} = \arg\max_v \prod_{j=1}^{5} p_j(x^j \mid y = v)p(y = v)$$

$$= \arg\max(\prod_{j=1}^{5} p_j(x^j \mid y = \text{Physics})p(y = \text{Physics}),$$

$$\prod_{j=1}^{5} p_j(x^j \mid y = \text{Biology})p(y = \text{Biology}),$$

$$\prod_{j=1}^{5} p_j(x^j \mid y = \text{Chemistry})p(y = \text{Chemistry}))$$

$$= \arg\max((0.1 * 0.005 * 0.95 * 0.001 * 0.005) * 0.35,$$

$$(0.01 * 0.03 * 0.999 * 0.1 * 0.006) * 0.40,$$

$$(0.2 * 0.05 * 0.95 * 0.008 * 0.003) * 0.25)$$

$$= \arg\max(8.3125 \cdot 10^{-10}, 7.1928 \cdot 10^{-8}, 5.7000 \cdot 10^{-8})$$

$$= \arg 7.1928 \cdot 10^{-8}$$

$$= \text{Biology}$$

Next, we classify B.

$$\hat{y} = \arg\max_v \prod_{j=1}^{5} p_j(x^j \mid y = v)p(y = v)$$

$$= \arg\max(\prod_{j=1}^{5} p_j(x^j \mid y = \text{Physics})p(y = \text{Physics}),$$

$$\prod_{j=1}^{5} p_j(x^j \mid y = \text{Biology})p(y = \text{Biology}),$$

$$\prod_{j=1}^{5} p_j(x^j \mid y = \text{Chemistry})p(y = \text{Chemistry}))$$

$$= \arg\max((0.1 * 0.005 * 0.05 * 0.999 * 0.995) * 0.35,$$

$$(0.01 * 0.03 * 0.001 * 0.9 * 0.994) * 0.40,$$

$$(0.2 * 0.05 * 0.05 * 0.992 * 0.997) * 0.25)$$

$$= \arg\max(8.6975 \cdot 10^{-6}, 1.0735 \cdot 10^{-7}, 1.2363 \cdot 10^{-4})$$

$$= \arg 1.2363 \cdot 10^{-4}$$

$$= \text{Chemistry}$$

Problem #5

Consider the following table of observations:

No.	Outlook	Temperature	Humidity	Windy	Play Golf?
1	sunny	hot	high	false	N
2	sunny	hot	high	true	N
3	overcast	hot	high	false	Y
4	rain	mild	high	false	Y
5	rain	cool	normal	false	Y
6	rain	cool	normal	true	N
7	overcast	cool	normal	true	Y
8	sunny	mild	high	false	N
9	sunny	cool	normal	false	Y
10	rain	mild	normal	false	Y
11	sunny	mild	normal	true	Y
12	overcast	mild	high	true	Y
13	overcast	hot	normal	false	Y
14	rain	mild	high	true	N

From the classified examples in the above table, construct two decision trees (by hand) for the classification "Play Golf." For the first tree, use Temperature as the root node. (This is a

really bad choice.) Continue the construction of tree as discussed in class for the subsequent nodes using information gain. Remember that different attributes can be used in different branches on a given level of the tree. For the second tree, follow the Decision Tree Learning algorithm described in class. At each step, choose the attribute with the highest information gain. Work out the computations of information gain by hand and draw the decision tree.

For construction of the decision trees, we follow the Decision Tree Learning Algorithm. For any given node, we pick the attribute with the largest information gain (IG). If I is the entropy function, then we always select the attribute v such that:

$$\text{Selected Attribute} = \max \arg_v \text{IG} = \max \arg_v \left(I_{\text{before}} - I_{\text{after}} \right)$$

Note that I_{after} is dependent upon v.

For the first tree, we are forced to use "Temperature" as the root node. As we shall see in construction of the second tree, "Temperature" is the worst of the four choices. Without considering the IG of the decision, we propagate examples into their respective branches. Four go to the 'hot' branch; six go to the 'mild' branch; four go to the 'cool' branch.

Before diving into the messier branches, I will first finish the 'hot' branch. If we choose "Outlook" as the next attribute to classify by, the IG is maximized. Two examples go to the 'sunny' branch (both 'N'), while the other two examples go to the 'overcast' branch (both 'Y'). No examples are propagated to the 'rain' branch, so its decision is based on the majority rule at the parent node ("Outlook" in this case). Since the split is even, the result is 'Undecided' or 'Error.'

Next, we will finish the 'mild' branch stemming from the root node. The best choice is Outlook ($\text{IG}_{\text{Outlook}} = 0.126 > \text{IG}_{\text{Humidity}} = \text{IG}_{\text{Windy}} = -0.082$). We note that only one instance was propagated to the 'overcast' branch, so it becomes a leaf node with decision 'Y'. Continuing down the 'sunny' branch, we note that the "Humidity" attribute maximizes the IG (i.e. all propagated instances belong to one class), resulting in the 'normal' branch leading to a 'Y' leaf node and the 'high' branch leading to an 'N' leaf node. Lastly, we continue down the 'mild'–'rain' branch. From here, it's easy to see that the "Windy" attribute fully segregates the classes, maximizing its IG. The 'true' branch leads to an 'N' leaf node, while 'false' leads to a 'Y' leaf node.

Finally, we finish the tree by completing the 'cool' branch stemming from the root node. The best attribute choice is either "Windy" or "Outlook" since they have identical IGs ($\text{IG}_{\text{Outlook}} = \text{IG}_{\text{Windy}} = 0.3113 > \text{IG}_{\text{Humidity}} = 0$). We will pick "Windy" since the result is essentially the same either way. The 'false' branch leads to a 'Y' leaf node, while the 'true' branch leads to our last two examples (one 'Y' and one 'N'). From here (node located at 'cool'–'true'), we select "Outlook" as our final attribute choice since it maximizes IG. One instance is propagated to the 'rain' branch (resulting in an 'N' leaf node), while the other is propagated to the 'overcast' branch (resulting in a 'Y' leaf node). The 'sunny' branch is determined by majority vote from the parent node. As before, the split is even, so the result is 'Undecided' or 'Error.'
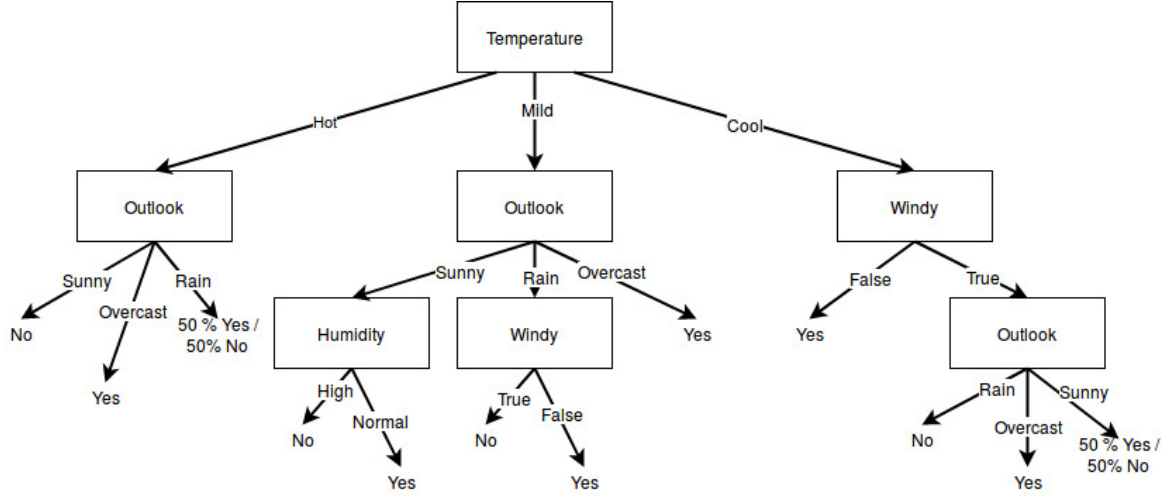
Finally, we construct the tree.

Figure 1: Decision Tree (poor root node choice)

For construction of the second tree, we follow the Decision Tree Learning Algorithm from the start. Beginning from the root node, we compute its attribute. The IG is computed in the following order: Outlook, Temperature, Humidity, Windy.

$$\text{Selected Attribute} = \max \arg_v \ (1 - (\frac{5}{14}I(\frac{3}{5},\frac{2}{5}) + \frac{4}{14}I(\frac{0}{4},\frac{4}{4}) + \frac{5}{14}I(\frac{2}{5},\frac{3}{5})),$$
$$1 - (\frac{4}{14}I(\frac{1}{2},\frac{1}{2}) + \frac{6}{14}I(\frac{2}{6},\frac{4}{6}) + \frac{4}{14}I(\frac{1}{4},\frac{3}{4})),$$
$$1 - (\frac{7}{14}I(\frac{3}{7},\frac{4}{7}) + \frac{7}{14}I(\frac{1}{7},\frac{6}{7})),$$
$$1 - (\frac{6}{14}I(\frac{3}{6},\frac{3}{6}) + \frac{8}{14}I(\frac{2}{8},\frac{6}{8})))$$
$$= \max \arg_v \ (0.3065, 0.0889, 0.2117, 0.1078)$$
$$= \text{Outlook}$$

As we can see, the root node will be "Outlook." In this case, we see that if the "Outlook" is 'overcast', the decision is always 'Y.' Thus, we obtain a leaf node there. For the other two branches– sunny and rain–, it's easy to see that maximum information gain can be obtained by using the "Humidity" attribute on the 'sunny' branch and the "Windy" attribute on the 'rain' branch. Because the information gain is maximized (i.e. all propagated instances belong to one class), there is no need to check other potential attributes. We instantly have leaf nodes. The resulting tree is below.
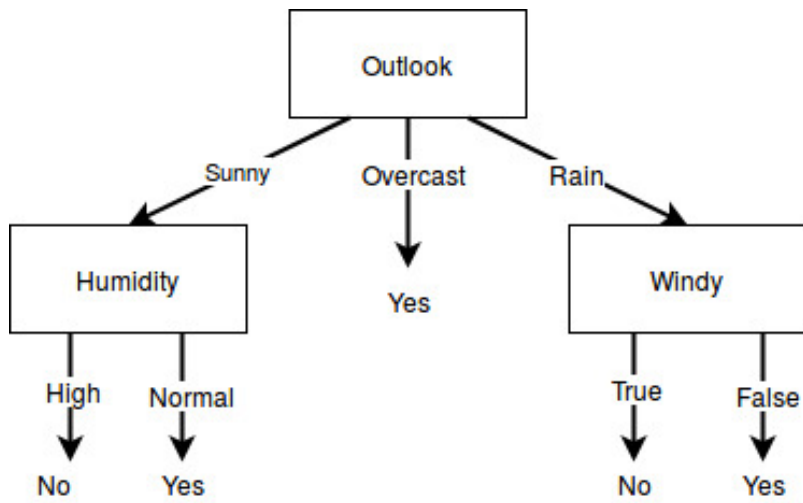
Figure 2: Decision Tree (good root node choice)

Note: all figures were drawn on "draw.io."